

introduction-to-regression-analysis

April 25, 2019

1 Regression analysis of medical data

Antti Honkela, antti.honkela@helsinki.fi

We have earlier done linear regression using scikit-learn. In this project work, however, we will use the [Statsmodels library](#). This is because Statsmodels has better statistical tools. In addition, it works better with Pandas' DataFrames, since it can specify the relation between the dependent and independent variables using a formula notation of column names of a DataFrame. Below is an example of a formula:

```
formula = "Y ~ X1 + X2"
```

So, the formula is given as a string where on the left side of the ~ character is the dependent variable, and on the right side the independent variables, separated using the + character. In this example the variable names Y, X1, and X2 refer to columns of a DataFrame.

```
In [3]: import numpy as np
import matplotlib.pyplot as plt
import statsmodels
import statsmodels.api as sm
import statsmodels.formula.api as smf
# plots a line given an intercept and a slope
from statsmodels.graphics.regressionplots import abline_plot
import pandas as pd
```

2 Multi-variable linear regression

Topics: - Multiple linear regression - Use of background variables to rectify regression - Interactions between variables - Choosing variables - Interpretation of estimation results

Multi-variable linear regression is used to model phenomena that depend on multiple variables. It can be used to adjust the model to consider confounding variables. It can also be used to recognize factors that have significant effect on a phenomenon.

Learning targets: - Fit multi-variable linear regression models in Python - Rectify regression models with background variables, and analyse the rectified models - Understand the principle of variable choosing - Understand most important restrictions of multiple linear regression models

Simple linear regression model is

$$y_i = \alpha + \beta x_i + \epsilon_i,$$

where

- y_i is the explained variable
- x_i is the explanatory variable
- β is the regression coefficient
- α is the constant term (intercept)
- ϵ_i is the residual.

Multi-variable linear regression model (or multiple liner regression model) is

$$y_i = \alpha + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \epsilon_i$$

- y_i is the explained variable
- x_{ij} are the explanatory variables $j = 1, \dots, p$
- β_j are the regression coefficients
- α is the constant term (intercept)
- ϵ_i is the residual.

The data can be represented as a design matrix that has variables as columns and observations as rows.

$$X = \begin{pmatrix} x_{11} & x_{12} & \cdots & x_{1p} \\ x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{np} \end{pmatrix}$$

The whole regression model in a matrix form is

$$y = \alpha \mathbf{1} + X\beta + \mathbf{ffl}$$

$$\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} = \alpha \begin{pmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{pmatrix} + \begin{pmatrix} x_{11} & x_{12} & \cdots & x_{1p} \\ x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \cdots & x_{np} \end{pmatrix} \begin{pmatrix} \beta_1 \\ \vdots \\ \beta_p \end{pmatrix} + \mathbf{ffl}$$

$$y_i = \alpha + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \epsilon_i$$

Or equivalently

$$y = (1 \ X) \begin{pmatrix} \alpha \\ \beta \end{pmatrix} + \mathbf{ffl}$$

$$\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} = \begin{pmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1p} \\ 1 & x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{np} \end{pmatrix} \begin{pmatrix} \alpha \\ \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{pmatrix} + \mathbf{ffl}$$

$$y_i = \alpha + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \epsilon_i$$

Or as Python expression:

```
y == np.concatenate([np.ones((len(x), 1)), X], axis=1) @ fit.params
```

2.0.1 An example using the Framingham Heart study

Data from the Framingham Heart study. In 1948, the study was initiated to identify the common factors or characteristics that contribute to CVD by following its development over time in group of participants who had not yet developed overt symptoms of CVD or suffered a heart attack or stroke. The researchers recruited 5,209 men and women between the ages of 30 and 62 from the town of Framingham, Massachusetts. Every two years, a series of extensive physical examinations and lifestyle interviews were conducted. This data set is subset of the Framingham Heart study data. The data is stored as 14 columns. Each row represents a single subject.

```
In [4]: # Load the data
```

```
fram = pd.read_csv('fram.txt', sep='\t')
fram.head()
```

```
Out[4]:
```

	ID	SEX	AGE	FRW	SBP	SBP10	DBP	CHOL	CIG	CHD	YRS_CHD	DEATH	\
0	4988	female	57	135	186	NaN	120	150	0	1	pre	7	
1	3001	female	60	123	165	NaN	100	167	25	0	16	10	
2	5079	female	54	115	140	NaN	90	213	5	0	8	8	
3	5162	female	52	102	170	NaN	104	280	15	0	10	7	
4	4672	female	45	99	185	NaN	105	326	20	0	8	10	

	YRS_DTH	CAUSE
0	11	unknown
1	17	unknown
2	13	unknown
3	11	unknown
4	17	unknown

ID	Explanation
SEX	Gender
AGE	Age at the start of the study
FRW	Weight in relation to groups median
SBP	Systolic Blood Pressure
DBP	Diastolic Blood Pressure
CHOL	Cholestherol level
CIG	Smoking (cigarets per day)

As an example, let's predict the systolic blood pressure using the weight.

```
In [5]: fit = smf.ols('SBP ~ FRW', data=fram).fit()
print(fit.summary())
```

```
OLS Regression Results
=====
Dep. Variable:          SBP      R-squared:          0.110
Model:                  OLS      Adj. R-squared:      0.110
Method:                  Least Squares      F-statistic:      172.5
```

```

Date:                Thu, 25 Apr 2019    Prob (F-statistic):        3.18e-37
Time:                22:08:31           Log-Likelihood:           -6542.3
No. Observations:    1394               AIC:                     1.309e+04
Df Residuals:        1392               BIC:                     1.310e+04
Df Model:            1
Covariance Type:     nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	92.8658	4.264	21.778	0.000	84.501	101.231
FRW	0.5241	0.040	13.132	0.000	0.446	0.602

```

Omnibus:                338.464    Durbin-Watson:                1.756
Prob(Omnibus):           0.000    Jarque-Bera (JB):            883.998
Skew:                    1.271    Prob(JB):                    1.10e-192
Kurtosis:                5.959    Cond. No.                    643.

```


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

In [6]: fram.plot.scatter("FRW", "SBP")
        #abline(results.params.Intercept, results.params.FRW, col="red")
        abline_plot(intercept=fit.params.Intercept, slope=fit.params.FRW,
                    ax=plt.gca(), color="red");

```



Next we rectify the model using background variables.

Assumptions of a regression model: 1. Relevance of data to the research question 2. Linearity and additivity 3. Independence of residuals 4. Constancy of variance of residuals 5. Normal distribution of residuals

Do these hold now?

In multiple-variable regression we add the background variables as explanators. Note: this rectification is linear and additive. In principle one should include all background variables, but estimation using too many variables can be unreliable.

Let's first consider a binary variable (gender).

```
In [7]: # Incorporate the gender
fit=smf.ols('SBP ~ FRW + SEX', data=fram).fit()
print(fit.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          SBP      R-squared:                0.118
Model:                  OLS      Adj. R-squared:           0.117
Method:                 Least Squares      F-statistic:        92.94
Date:                  Thu, 25 Apr 2019      Prob (F-statistic):    1.31e-38
Time:                  22:08:31      Log-Likelihood:       -6536.3
No. Observations:      1394      AIC:                  1.308e+04
Df Residuals:          1391      BIC:                  1.309e+04
Df Model:               2
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept              97.6608      4.466      21.866      0.000      88.899     106.422
SEX[T.male]           -4.9701      1.432      -3.470      0.001      -7.780      -2.161
FRW                     0.5010      0.040     12.431      0.000      0.422      0.580
=====
Omnibus:               330.964      Durbin-Watson:           1.768
Prob(Omnibus):         0.000      Jarque-Bera (JB):        851.938
Skew:                  1.249      Prob(JB):                1.01e-185
Kurtosis:              5.902      Cond. No.                 680.
=====
```


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Next we visualize men separately from women.

```
In [8]: fram.plot.scatter("FRW", "SBP")
int1 = fit.params.Intercept + fit.params["SEX[T.male]"]
int2 = fit.params.Intercept
slope=fit.params.FRW
```

```
abline_plot(intercept=int1, slope=slope, ax=plt.gca(), color="blue", label="male")
abline_plot(intercept=int2, slope=slope, ax=plt.gca(), color="red", label="female")
plt.legend();
```



The previous model acknowledged the gender in the intercept, but not in the slope. We improve the model by adding an *interaction term* `FRW:SEX`. Interaction is the product of the two variables. (Note that in these dependence formulas `A * B` is an abbreviation for `A + B + A:B`. The `*` character is not often used in the formulas.)

```
In [9]: # Include both gender and its interaction with the weight
fit2=smf.ols('SBP ~ FRW + SEX + FRW:SEX', data=fram).fit()
print(fit2.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          SBP      R-squared:                0.118
Model:                  OLS      Adj. R-squared:           0.116
Method:                 Least Squares      F-statistic:         61.92
Date:                  Thu, 25 Apr 2019      Prob (F-statistic):      1.42e-37
Time:                  22:08:32      Log-Likelihood:         -6536.3
No. Observations:      1394      AIC:                    1.308e+04
Df Residuals:          1390      BIC:                    1.310e+04
Df Model:               3
Covariance Type:       nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
=====
```

```

-----
Intercept          98.0027    5.324    18.408    0.000    87.559    108.446
SEX[T.male]       -6.0457    9.219    -0.656    0.512   -24.130    12.038
FRW                0.4979    0.048    10.288    0.000     0.403     0.593
FRW:SEX[T.male]    0.0103    0.087     0.118    0.906    -0.161     0.182
=====
Omnibus:                331.026  Durbin-Watson:                1.768
Prob(Omnibus):          0.000  Jarque-Bera (JB):            852.312
Skew:                   1.250  Prob(JB):                    8.37e-186
Kurtosis:               5.903  Cond. No.:                   1.66e+03
=====

```

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.66e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [11]: *# Renormalize to ease interpretation of the model parameters*

```

fram["sAGE"] = rescale(fram.AGE)
fram["sFRW"] = rescale(fram.FRW)
fram["sCHOL"] = rescale(fram.CHOL)
fram["sCIG"] = rescale(fram.CIG)
# Note: No need to scale the variable SEX

```

In [12]: *# Now with renormalized variables*

```

fit3=smf.ols('SBP ~ sFRW + SEX + sFRW:SEX', data=fram).fit()
print(fit3.summary())

```

OLS Regression Results

```

=====
Dep. Variable:          SBP    R-squared:                0.118
Model:                  OLS    Adj. R-squared:            0.116
Method:                 Least Squares    F-statistic:            61.92
Date:                  Thu, 25 Apr 2019    Prob (F-statistic):      1.42e-37
Time:                  22:08:32    Log-Likelihood:          -6536.3
No. Observations:      1394    AIC:                    1.308e+04
Df Residuals:          1390    BIC:                    1.310e+04
Df Model:               3
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	150.4587	0.984	152.940	0.000	148.529	152.389
SEX[T.male]	-4.9569	1.437	-3.449	0.001	-7.776	-2.138
sFRW	17.6762	1.718	10.288	0.000	14.306	21.047
sFRW:SEX[T.male]	0.3669	3.106	0.118	0.906	-5.727	6.461

```

=====

```


Omnibus:	331.026	Durbin-Watson:	1.768
Prob(Omnibus):	0.000	Jarque-Bera (JB):	852.312
Skew:	1.250	Prob(JB):	8.37e-186
Kurtosis:	5.903	Cond. No.	5.27

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [13]: p=fit3.params
         fram.plot.scatter("sFRW", "SBP")
         #abline(p.Intercept + p["SEX[T.male]"],
         #       p.sFRW + p["sFRW:SEX[T.male]"], col="blue", label="male")
         #abline(p.Intercept, p.sFRW, col="red", label="female")
         int1 = p.Intercept + p["SEX[T.male]"]
         int2 = p.Intercept
         slope1 = p.sFRW + p["sFRW:SEX[T.male]"]
         slope2 = p.sFRW
         abline_plot(intercept=int1, slope=slope1, ax=plt.gca(), color="blue", label="male")
         abline_plot(intercept=int2, slope=slope2, ax=plt.gca(), color="red", label="female")
         plt.legend();
```



2.1 Background variables

ID	Explanation
SEX	Gender
AGE	Age at the start of the study
FRW	Weight in relation to groups median
SBP	Systolic Blood Pressure
DBP	Diastolic Blood Pressure
CHOL	Cholestherol level
CIG	Smoking (cigarets per day)

Next we add a continuous background variable: cholesterol.

```
In [14]: fit4=smf.ols('SBP ~ sFRW + SEX + sFRW:SEX + sCHOL', data=fram).fit()
print(fit4.summary())
```

OLS Regression Results						
=====						
Dep. Variable:	SBP	R-squared:	0.125			
Model:	OLS	Adj. R-squared:	0.123			
Method:	Least Squares	F-statistic:	49.75			
Date:	Thu, 25 Apr 2019	Prob (F-statistic):	3.67e-39			
Time:	22:08:32	Log-Likelihood:	-6530.4			
No. Observations:	1394	AIC:	1.307e+04			
Df Residuals:	1389	BIC:	1.310e+04			
Df Model:	4					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	150.0166	0.988	151.776	0.000	148.078	151.955
SEX[T.male]	-4.0700	1.455	-2.798	0.005	-6.923	-1.216
sFRW	17.7594	1.712	10.375	0.000	14.402	21.117
sFRW:SEX[T.male]	-0.1272	3.098	-0.041	0.967	-6.204	5.950
sCHOL	4.9197	1.433	3.433	0.001	2.108	7.731
=====						
Omnibus:	327.586	Durbin-Watson:	1.774			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	843.566			
Skew:	1.237	Prob(JB):	6.64e-184			
Kurtosis:	5.899	Cond. No.	5.28			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [15]: # Add interactions between variables
fit4=smf.ols('SBP ~ sFRW + SEX + sFRW:SEX + sCHOL + sCHOL:sFRW + sCHOL:SEX',
            data=fram).fit()
print(fit4.summary())
```

OLS Regression Results

Dep. Variable:	SBP	R-squared:	0.127			
Model:	OLS	Adj. R-squared:	0.123			
Method:	Least Squares	F-statistic:	33.58			
Date:	Thu, 25 Apr 2019	Prob (F-statistic):	5.65e-38			
Time:	22:08:32	Log-Likelihood:	-6529.2			
No. Observations:	1394	AIC:	1.307e+04			
Df Residuals:	1387	BIC:	1.311e+04			
Df Model:	6					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]


Intercept	149.9420	0.994	150.850	0.000	147.992	151.892
SEX[T.male]	-4.0980	1.455	-2.816	0.005	-6.953	-1.243
sFRW	17.9750	1.718	10.466	0.000	14.606	21.344
sFRW:SEX[T.male]	0.2760	3.108	0.089	0.929	-5.821	6.373
sCHOL	5.5034	1.861	2.958	0.003	1.853	9.154
sCHOL:SEX[T.male]	-1.3225	2.939	-0.450	0.653	-7.087	4.442
sCHOL:sFRW	3.9052	2.741	1.425	0.154	-1.471	9.282
=====						
Omnibus:	318.099	Durbin-Watson:	1.769			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	798.422			
Skew:	1.212	Prob(JB):	4.22e-174			
Kurtosis:	5.805	Cond. No.	5.32			
=====						

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Normalized data (rescale) allows analysis of the importance of variables. An interpretation: how much does a change of 2*standard deviation affect the explained variable. In the following we visualize women with either low, medium or high cholesterol.

```
In [16]: p=fit4.params
         fram.plot.scatter("sFRW", "SBP")
         abline_plot(intercept=p.Intercept - p["sCHOL"], slope=p.sFRW - p["sCHOL:sFRW"],
                     ax=plt.gca(), color="blue", label="low")
         abline_plot(intercept=p.Intercept, slope=p.sFRW,
                     ax=plt.gca(), color="magenta", label="mid")
         abline_plot(intercept=p.Intercept + p["sCHOL"], slope=p.sFRW + p["sCHOL:sFRW"],
                     ax=plt.gca(), color="red", label="high")
         plt.legend();
```




Below is the same analysis but in separate visualizations.

```
In [17]: fig, ax = plt.subplots(1,3, subplot_kw={"xlim": (-1.6, 3.3), "ylim": (80,310),
                                                "xlabel": "sFRW", "ylabel": "SBP"},
                                figsize=(14, 4))
ax[0].scatter(fram.sFRW[(fram.SEX=="female") & (fram.sCHOL < -0.5)],
              fram.SBP[(fram.SEX=="female") & (fram.sCHOL < -0.5)])
abline_plot(p.Intercept - p["sCHOL"],
            p.sFRW - p["sCHOL:sFRW"], color="blue", label="low", ax=ax[0])
ax[0].set_title("female, low CHOL")

ax[1].scatter(fram.sFRW[(fram.SEX=="female") & (fram.sCHOL > -0.5) &
                        (fram.sCHOL < 0.5)],
              fram.SBP[(fram.SEX=="female") & (fram.sCHOL > -0.5) &
                        (fram.sCHOL < 0.5)])
abline_plot(p.Intercept, p.sFRW, color="magenta", label="mid", ax=ax[1])
ax[1].set_title("female, mid CHOL")

ax[2].scatter(fram.sFRW[(fram.SEX=="female") & (fram.sCHOL > 0.5)],
              fram.SBP[(fram.SEX=="female") & (fram.sCHOL > 0.5)])
abline_plot(p.Intercept + p["sCHOL"],
            p.sFRW + p["sCHOL:sFRW"], color="red", label="high", ax=ax[2])
ax[2].set_title("female, high CHOL")
```

```
Out[17]: Text(0.5, 1.0, 'female, high CHOL')
```



2.1.1 Prediction and generalization

Model's predictive accuracy in the data it was learned from does not give a good picture of its predictive capabilities: the model can be overfitted. A better estimate for the predictive accuracy can be obtained using cross validation: 1. Divide the data into parts for fitting and for validation 2. The model is fitted in a part of the data (training data) 3. The models is tested on another part of the data (test data). Then prediction error is computed. 4. This is repeated for a wanted number of divisions of the data

One model:

```
In [19]: train, test = train_test_split(fram)                # Split the data into two parts
         fit = smf.ols('SBP ~ sFRW + SEX + sCHOL', data=train).fit() # Fit the model
         pred = fit.predict(test)                               # Compute predictions
         rmse = np.sqrt(np.mean((pred - test.SBP)**2))         # Root mean square error
         rmse
```

```
Out[19]: 27.234790574539684
```

Another model:

```
In [20]: train, test = train_test_split(fram)
         fit = smf.ols('SBP ~ sFRW + SEX + sCHOL + sFRW:SEX + sCHOL:sFRW + sCHOL:SEX',
                        data=train).fit()
         pred = fit.predict(test)
         rmse = np.sqrt(np.mean((pred - test.SBP)**2))
         rmse
```

```
Out[20]: 26.560076340037707
```

Let's repeat this random data splitting 100 times for both models and compute the average RMSEs:

```
In [21]: error_basic=[]
         error_interact=[]
         np.random.seed(9)
```

```

for i in range(100):
    train, test = train_test_split(fram)
    fit1 = smf.ols('SBP ~ sFRW + SEX + sCHOL', data=train).fit()
    fit2 = smf.ols('SBP ~ sFRW + SEX + sCHOL + sFRW:SEX + sCHOL:sFRW + sCHOL:SEX',
                  data=train).fit()
    pred1 = fit1.predict(test)
    pred2 = fit2.predict(test)
    error_basic.append(np.sqrt(np.mean((pred1 - test.SBP)**2)))
    error_interact.append(np.sqrt(np.mean((pred2 - test.SBP)**2)))
pd.Series(error_basic).mean(), pd.Series(error_interact).mean()

```

Out [21]: (26.18486951586637, 26.259898342406064)

We can use the *Mann–Whitney U test* to see whether the prediction errors differ significantly.

```

In [22]: statsmodels.stats.stattools.stats.mannwhitneyu(error_basic, error_interact,
               alternative="two-sided")

```

Out [22]: MannwhitneyuResult(statistic=4857.0, pvalue=0.7277028167772348)

Let's change the first model and redo the experiment:

```

In [23]: error_basic=[]
         error_interact=[]
         np.random.seed(9)
         for i in range(100):
             train, test = train_test_split(fram)
             fit1 = smf.ols('SBP ~ sFRW ', data=train).fit()
             fit2 = smf.ols('SBP ~ sFRW + SEX + sCHOL + sFRW:SEX + sCHOL:sFRW + sCHOL:SEX',
                           data=train).fit()
             pred1 = fit1.predict(test)
             pred2 = fit2.predict(test)
             error_basic.append(np.sqrt(np.mean((pred1 - test.SBP)**2)))
             error_interact.append(np.sqrt(np.mean((pred2 - test.SBP)**2)))
         pd.Series(error_basic).mean(), pd.Series(error_interact).mean()

```

Out [23]: (26.38389912270515, 26.259898342406064)

```

In [24]: statsmodels.stats.stattools.stats.mannwhitneyu(error_basic, error_interact,
               alternative="two-sided")

```

Out [24]: MannwhitneyuResult(statistic=5222.0, pvalue=0.5883624055865182)

Now let's change the second model:

```

In [25]: error_basic=[]
         error_interact=[]
         np.random.seed(9)
         for i in range(100):
             train, test = train_test_split(fram)

```

```

fit1 = smf.ols('SBP ~ sFRW ', data=train).fit()
fit2 = smf.ols('SBP ~ sFRW + SEX + sCHOL + sAGE + sCIG', data=train).fit()
pred1 = fit1.predict(test)
pred2 = fit2.predict(test)
error_basic.append(np.sqrt(np.mean((pred1 - test.SBP)**2)))
error_interact.append(np.sqrt(np.mean((pred2 - test.SBP)**2)))
pd.Series(error_basic).mean(), pd.Series(error_interact).mean()

```

Out [25]: (26.38389912270515, 25.924525100203073)

```

In [26]: statsmodels.stats.stattools.stats.mannwhitneyu(error_basic, error_interact,
               alternative="two-sided")

```

Out [26]: MannwhitneyuResult(statistic=5742.0, pvalue=0.0700213003206784)

3 Logistic regression

Topics:

- Logistic regression model
- Classification problems

Learning targets:

- Fit logistic regression models with Python
- Interpret the estimated regression models

3.1 Regression model is transformations of variables

Multi-variable linear regression model:

$$y_i = \alpha + \beta_1 x_{i1} + \dots + \beta_p x_{ip} + \epsilon_i$$

The model is very flexible with respect to the variables x_{ij} and y_i : the variables need not be direct observations, for example the interaction terms like $SEX:sWHP$. Also transformations of variables are permitted, for example $SBP \sim \log(FRW) + sFRW + SEX + SEX:sFRW$.

For example, logarithm transform is often useful for variables, whose range is large and whose effect can be expected to saturate. An example: $\log(SBP) \sim \log(FRW) + SEX + SEX:\log(FRW)$

$$SBP = \alpha FRW^{\beta_1} \exp(SEX)^{\beta_2} FRW^{SEX \cdot \beta_3}$$

3.2 Binary target variable (classification)


It is not sensible to try to predict a binary variable directly using linear regression. In general, we want to predict $p(y_i = \text{TRUE} \mid X)$. In linear regression the possible values are in the interval $(-\infty, \infty)$, whereas probabilities are in the interval $[0, 1]$. The idea is to transform the unrestricted predictions to probabilities.

```

In [27]: def logistic(x):
         return 1.0 / (1.0 + np.exp(-x))

```

```
In [28]: X=np.linspace(-8, 8, 100)
plt.plot(X, logistic(X));
```



$$\text{logit}^{-1}(x) = \frac{1}{1 + \exp(-x)}$$

Logistic transform is non-linear: same change in input produces different changes in probabilities. The speed of change is at its largest at the point $x = 0 : f'(0) = 1/4$. Logistic regression is the most common tool for classification. It can also be used to recognize variables that are important to the classification.

Let's continue with the `fram` data. First we define a diagnose for high blood pressure.

```
In [29]: fram["HIGH_BP"] = (fram.SBP >= 140) | (fram.DBP >= 90)
fram.HIGH_BP.head()
```

```
Out[29]: 0    True
         1    True
         2    True
         3    True
         4    True
         Name: HIGH_BP, dtype: bool
```

```
In [30]: fram.HIGH_BP.value_counts()
```

```
Out[30]: True      906
         False    488
         Name: HIGH_BP, dtype: int64
```

```
In [31]: fram.HIGH_BP = fram.HIGH_BP.map(int)
```

Note that for boolean variables we use type `int` here instead of `bool`, because we want to make the encoding of booleans as integers explicit: 0 for False and 1 for True. (The implicit encoding of booleans as integers in statsmodels library is unfortunately inconsistent.)

```
In [32]: fram.HIGH_BP.mean()           # Fraction of observations with this diagnose
```

```
Out[32]: 0.6499282639885222
```

```
In [33]: fram.head()
```

```
Out[33]:
```

	ID	SEX	AGE	FRW	SBP	SBP10	DBP	CHOL	CIG	CHD	YRS_CHD	DEATH	\
0	4988	female	57	135	186	NaN	120	150	0	1	pre	7	
1	3001	female	60	123	165	NaN	100	167	25	0	16	10	
2	5079	female	54	115	140	NaN	90	213	5	0	8	8	
3	5162	female	52	102	170	NaN	104	280	15	0	10	7	
4	4672	female	45	99	185	NaN	105	326	20	0	8	10	

	YRS_DTH	CAUSE	sAGE	sFRW	sCHOL	sCIG	HIGH_BP
0	11	unknown	0.477764	0.834668	-0.914016	-0.346569	1
1	17	unknown	0.791473	0.496687	-0.730446	0.732493	1
2	13	unknown	0.164056	0.271367	-0.233727	-0.130757	1
3	11	unknown	-0.045083	-0.094779	0.489755	0.300868	1
4	17	unknown	-0.777070	-0.179274	0.986475	0.516680	1

Let's fit a logistic regression model:

```
In [34]: fit1 = smf.glm(formula="HIGH_BP ~ FRW", data=fram,
                        family=sm.families.Binomial(statsmodels.genmod.families.links.logit)).fit()
fit1.summary()
```

```
Out[34]: <class 'statsmodels.iolib.summary.Summary'>
"""
                        Generalized Linear Model Regression Results
=====
Dep. Variable:          HIGH_BP      No. Observations:          1394
Model:                  GLM          Df Residuals:              1392
Model Family:           Binomial     Df Model:                  1
Link Function:          logit        Scale:                    1.0000
Method:                 IRLS         Log-Likelihood:          -858.64
Date:                   Thu, 25 Apr 2019    Deviance:                1717.3
Time:                   22:08:40          Pearson chi2:            1.39e+03
No. Iterations:         4              Covariance Type:        nonrobust
=====
                        coef      std err          z      P>|z|      [0.025      0.975]
-----
Intercept             -2.8912      0.404      -7.158      0.000      -3.683      -2.100
FRW                    0.0339      0.004       8.650      0.000       0.026       0.042
=====
"""
```




```
In [35]: fit1.params
```

```
Out[35]: Intercept    -2.891166  
FRW                0.033852  
dtype: float64
```

The R^2 is not sensible now. Instead, we use *deviance*, which measures the error. Smaller value is better. The coefficients are mostly like in linear regression. Also, the significance interpretation is the same. Coefficient β : change of one unit in a variable causes a change in the probability which is at most $\beta/4$.

```
In [36]: # Visualization of the model  
plt.scatter(fram.FRW, fram.HIGH_BP, marker="d")  
X=np.linspace(40, 235, 100)  
plt.plot(X, logistic(X*fit1.params.FRW + fit1.params.Intercept))  
plt.xlabel("FRW")  
plt.ylabel("HIGH_BP")
```

```
Out[36]: Text(0, 0.5, 'HIGH_BP')
```



Next we add the gender and its interaction to the model:

```
In [37]: fit2 = smf.glm(formula="HIGH_BP ~ sFRW + SEX + SEX:sFRW", data=fram,  
                        family=sm.families.Binomial()).fit()  
fit2.summary()
```

```
Out [37]: <class 'statsmodels.iolib.summary.Summary'>
```


```
"""
```

Generalized Linear Model Regression Results


=====						
Dep. Variable:	HIGH_BP	No. Observations:	1394			
Model:	GLM	Df Residuals:	1390			
Model Family:	Binomial	Df Model:	3			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-856.87			
Date:	Thu, 25 Apr 2019	Deviance:	1713.7			
Time:	22:08:40	Pearson chi2:	1.39e+03			
No. Iterations:	4	Covariance Type:	nonrobust			
=====						
	coef	std err	z	P> z	[0.025	0.975]

Intercept	0.7631	0.082	9.266	0.000	0.602	0.925
SEX[T.male]	-0.1624	0.120	-1.350	0.177	-0.398	0.073
sFRW	1.0738	0.176	6.095	0.000	0.728	1.419
SEX[T.male]:sFRW	0.2709	0.287	0.943	0.346	-0.292	0.834
=====						
"""						

```
In [38]: plt.scatter(fram.sFRW, fram.HIGH_BP, marker="d")
X=np.linspace(-2, 4, 100)
p = fit2.params
plt.plot(X, logistic(X*p.sFRW + p.Intercept), color="red", label="female")
plt.plot(X, logistic(X*(p.sFRW + p["SEX[T.male]:sFRW"]) +
            p["SEX[T.male]"] + p.Intercept), color="blue", label="male")
plt.xlabel("Weight")
plt.ylabel("Pr(Has high BP)")
plt.legend();
```



```
In [39]: # We add a bit of random jitter to the y value
plt.scatter(fram.sFRW, fram.HIGH_BP + np.random.uniform(-0.05, 0.05, len(fram)),
            marker="d")
X=np.linspace(-2, 4, 100)
p = fit2.params
plt.plot(X, logistic(X*p.sFRW + p.Intercept), color="red", label="female")
plt.plot(X, logistic(X*(p.sFRW + p["SEX[T.male]:sFRW"]) +
                    p["SEX[T.male]"] + p.Intercept), color="blue", label="male")
plt.xlabel("Weight")
plt.ylabel("Pr(Has high BP)")
plt.legend();
```



3.3 Prediction

The `fittedvalues` attribute contains the predicted probabilities for each data point. Let's compute the fraction of mispredictions:

```
In [40]: error_rate = np.mean(((fit2.fittedvalues < 0.5) & fram.HIGH_BP) |
                               ((fit2.fittedvalues > 0.5) & ~fram.HIGH_BP))
        error_rate
```

```
Out[40]: 0.35581061692969873
```

What is a good error rate? For a random guess (tossing a coin): error rate is 50%. If we always choose the most common option, the error rate is less than 50%, sometimes a lot less.

```
In [41]: print("Base rate:", 1-np.mean(fram.HIGH_BP))
```

```
Base rate: 0.3500717360114778
```

3.4 Cross validation

```
In [42]: train, test = train_test_split(fram, seed=0)
        print(len(train), len(test))
        fit = smf.glm(formula="HIGH_BP ~ sFRW + SEX + SEX:sFRW", data=train,
                       family=sm.families.Binomial(statsmodels.genmod.families.links.logit))
```

```

print(fit.summary())
#print(test.head())
pred = fit.predict(test, transform=True)
#print(pred.describe())
#print("Min:", pred.min())
#print("Max:", pred.max())
error_rate = np.mean(((pred < 0.5) & (test.HIGH_BP==1)) |
                      ((pred > 0.5) & (test.HIGH_BP==0)))
print(error_rate, 1 - test.HIGH_BP.mean())

```

1115 279

Generalized Linear Model Regression Results

```

=====
Dep. Variable:          HIGH_BP    No. Observations:          1115
Model:                  GLM        Df Residuals:                1111
Model Family:           Binomial   Df Model:                    3
Link Function:           logit     Scale:                     1.0000
Method:                  IRLS      Log-Likelihood:           -689.76
Date:                    Thu, 25 Apr 2019    Deviance:                 1379.5
Time:                    22:08:41    Pearson chi2:              1.11e+03
No. Iterations:          4          Covariance Type:          nonrobust
=====

```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	0.8008	0.092	8.752	0.000	0.621	0.980
SEX[T.male]	-0.2329	0.134	-1.743	0.081	-0.495	0.029
sFRW	0.9641	0.197	4.895	0.000	0.578	1.350
SEX[T.male]:sFRW	0.1760	0.315	0.558	0.577	-0.442	0.794

```

=====

```

0.35125448028673834 0.36200716845878134

```

In [43]: error_model=[]
         error_null=[]
         np.random.seed(1)
         for i in range(100):
             train, test = train_test_split(fram)
             fit = smf.glm(formula="HIGH_BP ~ sFRW + SEX + SEX:sFRW", data=train,
                           family=sm.families.Binomial(statsmodels.genmod.families.links.logit))
             #print(model.summary())
             pred = fit.predict(test, transform=True)
             error_rate = np.mean(((pred < 0.5) & (test.HIGH_BP==1)) |
                                   ((pred > 0.5) & (test.HIGH_BP==0)))
             error_model.append(error_rate)
             error_null.append((1 - test.HIGH_BP).mean())
         #for model, null in zip(error_model, error_null):
         #    print(model, null)
         pd.Series(error_model).mean(), pd.Series(error_null).mean()

```

Out [43]: (0.3520071684587814, 0.3488530465949821)

Test to see whether the results are significantly different:

```
In [44]: statsmodels.stats.stattools.stats.mannwhitneyu(error_model, error_null,
               alternative="two-sided")
```

Out [44]: MannwhitneyuResult(statistic=5312.5, pvalue=0.44539961439180353)

Let's define another diagnose:

```
In [45]: fram["HIGH_BP2"] = (fram.SBP > 140) | (fram.DBP > 90)
        fram["HIGH_BP2"] = fram["HIGH_BP2"].map(int)
        fram["HIGH_BP2"].mean()
```

Out [45]: 0.56025824964132

```
In [46]: error_model=[]
        error_null=[]
        np.random.seed(9)
        for i in range(100):
            train, test = train_test_split(fram)
            fit = smf.glm(formula="HIGH_BP2 ~ sFRW + SEX + SEX:sFRW", data=train,
                          family=sm.families.Binomial()).fit()
            #print(model.summary())
            pred = fit.predict(test)
            error_rate = np.mean(((pred < 0.5) & (test.HIGH_BP2==1)) |
                                ((pred > 0.5) & (test.HIGH_BP2==0)))
            error_model.append(error_rate)
            error_null.append((1-test.HIGH_BP2).mean())
        #for model, null in zip(error_model, error_null):
        #    print(model, null)
        pd.Series(error_model).mean(), pd.Series(error_null).mean()
```

Out [46]: (0.3860931899641578, 0.4403942652329749)

Test again:

```
In [47]: statsmodels.stats.stattools.stats.mannwhitneyu(error_model, error_null,
               alternative="two-sided")
```

Out [47]: MannwhitneyuResult(statistic=525.0, pvalue=7.3782255350127615e-28)